

FSST string compression

FSST = Fast Static Symbol Table

Peter Boncz (CWI)

Viktor Leis (FSU Jena)

Thomas Neumann (TU Munich)


<https://github.com/cwida/fsst>

String Compression in a DBMS

- Dictionary Compression
 - Whole string becomes 1 code, points into a dictionary D
 - works well if there are (relatively) few unique strings
- Heavy-weight/general-purpose Compression
 - Lempel-Zif plus possibly entropy coding
 - Zip, gzip, snappy, **LZ4**, zstd, ...
 - Block-based decompression

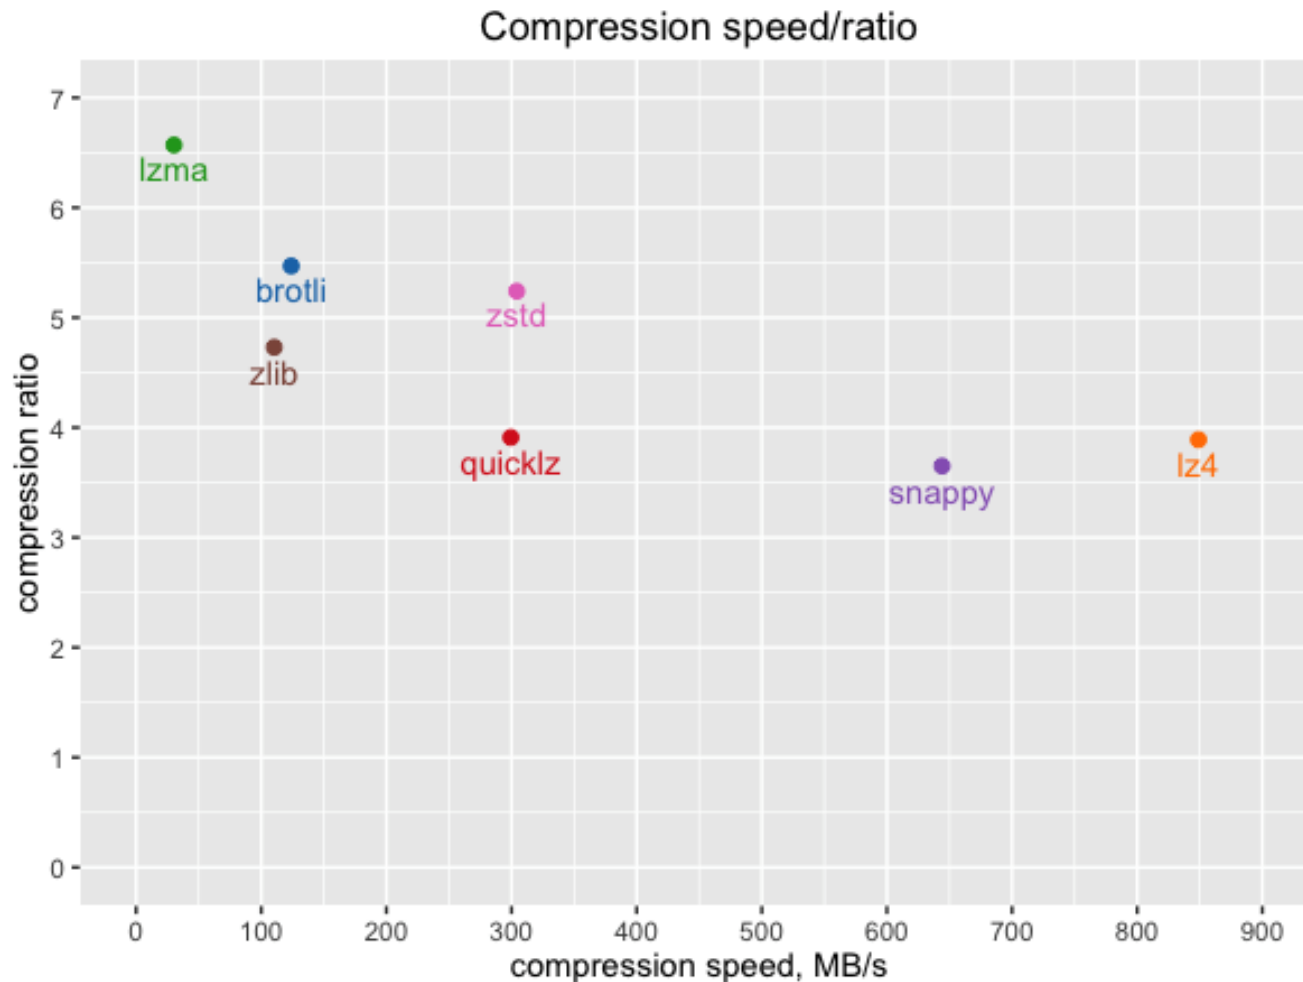
String Compression in a DBMS

www.percona.com/blog/2016/04/13/evaluating-database-compression-methods-update/



String Compression in a DBMS

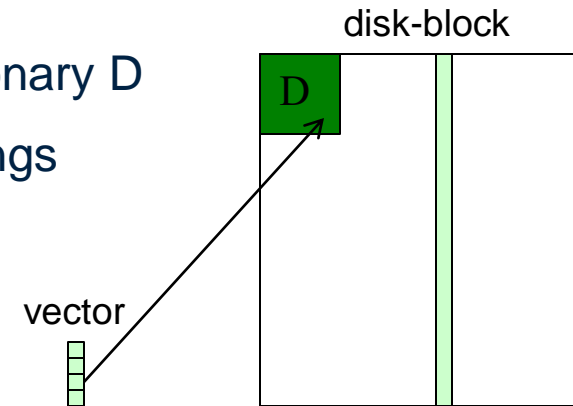
www.percona.com/blog/2016/04/13/evaluating-database-compression-methods-update/



“LZ4 is the best compression library for database applications, if high speed is desired”

String Compression in a DBMS


- Dictionary Compression
 - Whole string becomes 1 code, points into a dictionary D
 - works well if there are (relatively) few unique strings




- Heavy-weight/general-purpose Compression
 - Lempel-Zipf plus possibly entropy coding
 - Zip, gzip, snappy, **LZ4**, zstd, ...
 - Block-based decompression

String Compression in a DBMS

- Dictionary Compression
 - Whole string becomes 1 code, points into a dictionary D
 - works well if there are (relatively) few unique strings



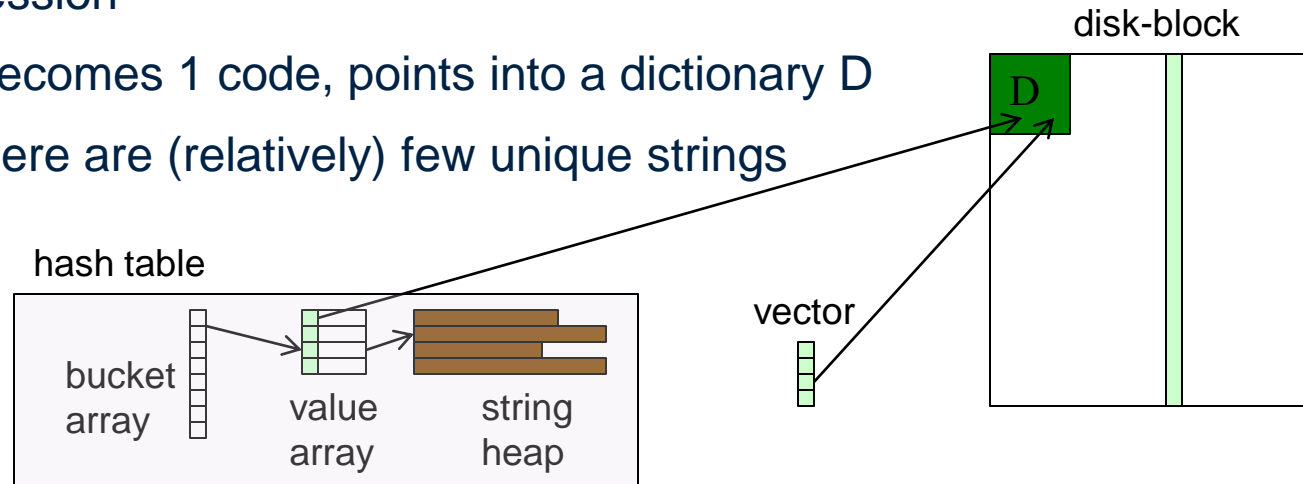
- Heavy-weight/general-purpose Compression
 - Lempel-Zipf plus possibly entropy coding
 - Zip, gzip, snappy, **LZ4**, zstd, ...
 - Block-based decompression



String Compression in a DBMS


- Dictionary Compression

- Whole string becomes 1 code, points into a dictionary D
- works well if there are (relatively) few unique strings



- Heavy-weight/general-purpose Compression


- Lempel-Zipf plus possibly entropy coding
- Zip, gzip, snappy, **LZ4**, zstd, ...
- Block-based decompression



String Compression in a DBMS


- Dictionary Compression

- Whole string becomes 1 code, points into a dictionary D
- works well if there are (relatively) few unique strings



- Heavy-weight/general-purpose Compression

- Lempel-Zipf plus possibly entropy coding
- Zip, gzip, snappy, **LZ4**, zstd, ...
- Block-based decompression



- **must decompress (all=) unneeded values in scan**
- **cannot be leveraged in hash tables, sorting, network shuffles**
- **FSST targets compression of many small textual strings**

The Idea

- Encode strings as a sequence of bytes, where each byte $[0,254]$ is a

– **CODE**

- Each code stands for a 1-8 byte

– **SYMBOL**

corpus
(uncompressed)

```
http://in.tum.de
http://cwi.nl
www.uni-jena.de
www.wikipedia.org
http://www.vldb.org
...
```

symbol table

0	http://	7
1	www.	4
2	uni-jena	8
3	.de	3
4	.org	4
5	a	1
6	in.tum	6
7	cwi.nl	6
8	wikipedi	8
9	vldb	4
...		
255		

symbol length

corpus
(compressed)

```
063
07
123
1854
0194
...
```

- Byte 255 is special code marking

– **EXCEPTION**

followed by 1 uncompressed byte

Small symbol table(s):
RAM: 2.2KB,
disk/network: ~500B

Closest existing scheme is **RePair**, but is $>100x$ slower than FSST (both ways)

FSST Decoding

Algorithm 1 FSST-decoding

```
void decode(uint8_t*& in, uint8_t*& out,  
            uint64_t sym[255], uint8_t len[255]) {  
    uint8_t code = *in++;  
    if (code != 255) {  
        *((uint64_t*)out) = sym[code];  
        out += len[code];  
    } else { // escape code  
        *out++ = *in++;  
    }  
}
```

Fast FFST Decoding

- Fast-skip escapes, handle escapes with Duff's device

```
while (posOut+32 <= size && posIn+4 <= lenIn) {
    unsigned int nextBlock = *((unsigned int*) (strIn+posIn));
    unsigned int escapeMask = (nextBlock&0x80808080u)&((((~nextBlock)&0x7F7F7F7Fu)+0x7F7F7F7Fu)^0x80808080u);
    if (escapeMask == 0) {
        code = strIn[posIn++]; *(unsigned long*) (strOut+posOut) = symbol[code]; posOut += len[code];
        code = strIn[posIn++]; *(unsigned long*) (strOut+posOut) = symbol[code]; posOut += len[code];
        code = strIn[posIn++]; *(unsigned long*) (strOut+posOut) = symbol[code]; posOut += len[code];
        code = strIn[posIn++]; *(unsigned long*) (strOut+posOut) = symbol[code]; posOut += len[code];
    } else {
        unsigned firstEscapePos=__builtin_ctzl(escapeMask)>>3;
        switch(firstEscapePos) {
            case 3: code = strIn[posIn++]; *(unsigned long*) (strOut+posOut) = symbol[code]; posOut += len[code];
            case 2: code = strIn[posIn++]; *(unsigned long*) (strOut+posOut) = symbol[code]; posOut += len[code];
            case 1: code = strIn[posIn++]; *(unsigned long*) (strOut+posOut) = symbol[code]; posOut += len[code];
            case 0: posIn+=2; strOut[posOut++] = strIn[posIn-1]; /* decompress an escaped byte */
        }
    }
}
```

FSST encoding


Algorithm 2 FSST-encoding, given a symbol table.

```
void encode(uint8_t*& in, uint8_t*& out, SymbolTable& st) {  
    uint16_t pos = st.findLongestSymbol(in);  
    if (pos <= 255) { // no (real) symbol found  
        *(out++) = 255;  
        *(out++) = *(in++);  
    } else {  
        *(out++) = (uint8_t) pos;  
        in += st.symbols[pos].len; // symbol length in bytes  
    }  
}
```

FSST Compression in a DBMS


- Dictionary Compression

- column (green) contains dictionary codes
- dictionary can be FSST compressed (smaller)
- no decompression on scan; random access possible



- FSST Compression

- column (brown) contains offsets in string segment
- eg 64KB block (self-aligned in RAM) starts with symbol table
- vectors contain pointers into block (can infer st from pointer)
- no decompression on scan; random access possible



Challenge: Finding a Good Symbol Table

- Why is this hard? **Dependency Problem!**
- First attempt:
 - Put the corpus in a **suffix array**
 - Identify the 255 common substrings with most **gain** ($=\text{length} \times \text{frequency}$)
 - **Problem 1:**
 - Valuable symbols will be **overlapping** (they are not as valuable as they seem)
 - We tried compensating for overlap → did not work
 - **Problem 2:** (**greedy** encoding)
 - The encoding will not arrive at the start of the valuable symbol, because the previous encoded symbol ate away the first byte(s)
 - We tried dynamic programming encoding (slow!!) → no improvements

FSST bottom-up symbol table construction

- Evolutionary-style algorithm
- Starts with empty symbol table, uses 5 iterations:
 - We encode (a sample of) the plaintext with the current symbol table
 - We count the occurrence of each symbol
 - We count the occurrence of each two subsequent symbols
 - We also count single byte(-extension) frequencies, even if these are not symbols
 - Two subsequent symbols (or byte-extensions) generate a new concatenated symbol
 - We compute the gain ($\text{length} \times \text{freq}$) of all bytes, old symbols and concatenated symbols and insert the 255 best in the new symbol table

Algorithm Example ($\text{maxlen}=3$, $|\text{st}|=5$)

Uncompressed

t	u	m	c	w	i	t	u	m	v	l	d	b
---	---	---	---	---	---	---	---	---	---	---	---	---

 len = 13

empty symbol table

Algorithm Example (maxlen=3, |st|=5)

Uncompressed

t u m c w i t u m v l d b len = 13

empty symbol table

Iteration 1

\$ t \$ u \$ m \$ c \$ w \$... \$ d \$ b len = 26

Symbol table

symbol	um	tu	wi	cw	mc
len	2	2	2	2	2
count	2	2	1	1	1
gain	4	4	2	2	2

(len * cnt)

Algorithm Example (maxlen=3, |st|=5)

Iteration 1

\$ t \$ u \$ m \$ c \$ w \$... \$ d \$ b

len = 26

Symbol table

symbol	um	tu	wi	cw	mc
len	2	2	2	2	2
count	2	2	1	1	1
gain	4	4	2	2	2

(len * cnt)

Iteration 2

tu mc wi tu \$ m \$ v \$ l \$ d \$ b len = 14

Symbol table

symbol	tum	tu	wit	mcw	vl
len	3	2	3	3	2
count	2	2	1	1	1
gain	6	4	3	3	2

Algorithm Example ($\text{maxlen}=3$, $|\text{st}|=5$)

Iteration 2

tu mc wi tu \$ m \$ v \$ l \$ d \$ b len = 14

Symbol table

symbol	tum	tu	wit	mcw	vl
len	3	2	3	3	2
count	2	2	1	1	1
gain	6	4	3	3	2

Iteration 3

tum \$ c wit \$ u \$ m vl \$ d \$ b len = 13

Symbol table

symbol	mvl	cwi	vld	tum	wit
len	3	3	3	3	3
count	1	1	1	1	1
gain	3	3	3	3	3

Algorithm Example (maxlen=3, |st|=5)

Iteration 3

tum \$ c wit \$ u \$ m vl \$ d \$ b len = 13

Symbol table	symbol	mvl	cwi	vld	tum	wit
	len	3	3	3	3	3
	count	1	1	1	1	1
	gain	3	3	3	3	3

Iteration 4

tum cwi tum vld \$ b len = 6

Symbol table	symbol	tum	cwi	vld	b
	len	3	3	3	1
	count	2	1	1	1
	gain	6	3	3	1

Algorithm Example (maxlen=3, |st|=5)

Iteration 4

tum cwi tum vld \$ b len = 6

Symbol table

symbol	tum	cwi	vld	b
len	3	3	3	1
count	2	1	1	1
gain	6	3	3	1

Compressed

tum cwi tum vld b len = 5

Uncompressed

t u m c w i t u m v l d b len = 13

Making FSST encoding fast

- **findLongestSymbol()**
 - Finds the next symbol
 - How? Range-scan in sorted list, indexed by first byte
 - A for-loop
- Goal: encoding without **for-loop** and without **if-then-else**
- Idea: **Lossy Perfect Hash Table**
 - Perfect: no collisions. How? Throw away colliding symbol with **least gain**
 - Lossy, therefore. But: we keep filling it with candidates until full anyway
 - Hash table key is **next 3 bytes**
 - Use a **shortCodes [65536]** direct lookup array for the **next two bytes**
 - Append a **terminator** single-byte symbol to plaintext (typically byte 0)

Making FSST encoding fast

- Idea: **Lossy Perfect Hash Table**
 - Perfect: no collisions. How?
 - Throw away colliding symbol with **least gain**
 - Lossy, therefore.
 - But: we keep filling it with candidates until full anyway
 - Hash table key is **next 3 bytes**
 - Use a **shortCodes [65536]** direct lookup array for the **next two bytes**
 - Append a **terminator** single-byte symbol to plaintext (typically byte 0)

```

u16 shortCodes[65536]; // code:12,length:4
struct Symbol {
    union { char str[maxLength]; u64 num; } val;
    u32 icl; // icl = u64 ignoredBits:16,code:12,length:4
}
Symbol hashTab[4096];

```

+-----+-----+		
code	len	
+-----+-----+		

+-----+-----+-----+		
ignoredBits	code	len
+-----+-----+-----+		

Lossy Perfect Hash Table

- String encoding without **for-loop** and without **if-then-else**

```

auto hashFind = [&]() {
    out[1] = word; // dirty trick: speculatively write out escaped byte
    uint64_t idx = hash(word & 0xFFFFFFFF) & (st.hashTabSize-1);
    s = st.hashTab[idx]; // fetch symbol from hash table
    uint64_t clean = (-1LL >> /*ignoredBits*/ (uint8_t) s.icl);
    return (s.icl < 0xFFFFFFFF && s.val.num == (word & clean));
};

auto hashedKernel = [&]() {
    code = hashFind() ? (s.icl>>16) : code; // conditional move
    *out = code; // write out code byte (or 255)
    out += 1+((code>>8)&1); // increase with 1 or 2 (escape = 9th bit)
    cur += (code>>12); // symbol length is in bits [12..15] of code
}
);

void
encodeHash(uint8_t*& cur, uint8_t*& out, SymbolTable& st){
    Symbol s;
    uint64_t word = *(uint64_t*)cur;
    uint64_t code = st.shortCodes[word & 0xFFFF];
    hashedKernel();
}

```


AVX512 Implementation

- Idea: compress 8 strings in parallel (8 lanes of 64-bits)
 - $*3 = 24$ in parallel (unrolled loop) – not $*4$ because of register pressure
 - **job** queue: 511 byte (max) string chunks
 - Add terminator symbol to each chunk
 - Sort jobs on string length (longest first) – load balancing, keep lanes busy
 - 512 jobs of 511B input, 1024B output (768KB buffer)
 - Each iteration:
 - Insert new jobs in (any) free lanes (**expand-load**)
 - **findLongestSymbol()** in AVX512
 - Match 1 symbol in input, add 1 code to output strings (in each lane)
 - Involves 3x**gather** (2x **hashTab** 1x**shortCodes**) + 1x**scatter** (output)
 - Append finished jobs in result job array (**compress-store**)

Evaluation: dbtext corpus

- **machine-readable identifiers** (hex, yago, email, wiki, `uuid, urls2, urls),
- **human-readable names** (firstname, lastname, city, credentials, street, movies),
- **text** (faust, hamlet, chinese, japanese, wikipedia),
- **domain-specific codes** (genome, location)
- **TPC-H data** (c_name, l_comment, ps_comment)

name	avg len	example string	LZ4 factor	FSST factor
hex	8	DD5AF484	1.14×	2.11×
yago	19	Ralph_A._Brown	1.25×	1.63×
email	22	xnj_14@hotmail.com	1.55×	2.13×
wiki	23	Benzil	1.31×	1.63×
uuid	37	84e22ac0-2da5-11e8-9d15- ...	1.55×	2.44×
urls2	55	http://fr.wikipedia.org/ ...	1.75×	2.05×
urls	63	http://reference.data.go ...	2.77×	2.42×
firstname	7	RUSSEL	1.25×	2.04×
lastname	10	BALONIER	1.28×	1.97×
city	10	ROELAND PARK	1.37×	2.14×
credentials	11	PHD, HSPP	1.48×	2.31×
street	13	PURITAN AVENUE	1.60×	2.35×
movies	21	Return to 'Giant'	1.23×	1.66×
faust	24	Erleuchte mein bedÄijrftig Herz.	1.48×	1.87×
hamlet	30	<LINE>That to Laertes ...	2.13×	2.41×
chinese	87	道人决心消除肉会 ...	1.40×	1.69×
japanese	90	せん。しかし、...	1.84×	2.00×
wikipedia	130	Weniger hÄd'ufig fressen sie ...	1.45×	1.81×
genome	10	atagtgaag	1.59×	3.32×
location	40	(40.84242764486843, -73 ...	1.58×	2.51×
c_name	19	Customer#000010485	3.08×	3.80×
l_comment	27	nal braids nag carefully expres	2.22×	2.90×
ps_comment	124	c foxes. fluffily ironic ...	2.79×	3.40×

Note: traditional compression datasets (e.g. Silesia) contain >50% binary files. Our new corpus is representative for DB text.

FSST vs LZ4

- Note **first bar** with overall average (AVG)
- FSST has **better compression factor** and **better compression speed** than LZ4
— **equal decompression speed**




Figure 3: Relative performance of FSST versus LZ4 in terms of compression factor, compression speed, and decompression speed. Each data set is treated as a 8MB file.

LZ4 as a database compressor

- It does not make sense to use LZ4 to compress strings one-by-one (“line”), even when using a pretrained zstd dictionary (“dict”). It is slow and has bad compression factor. General-purpose compression should be block-based.

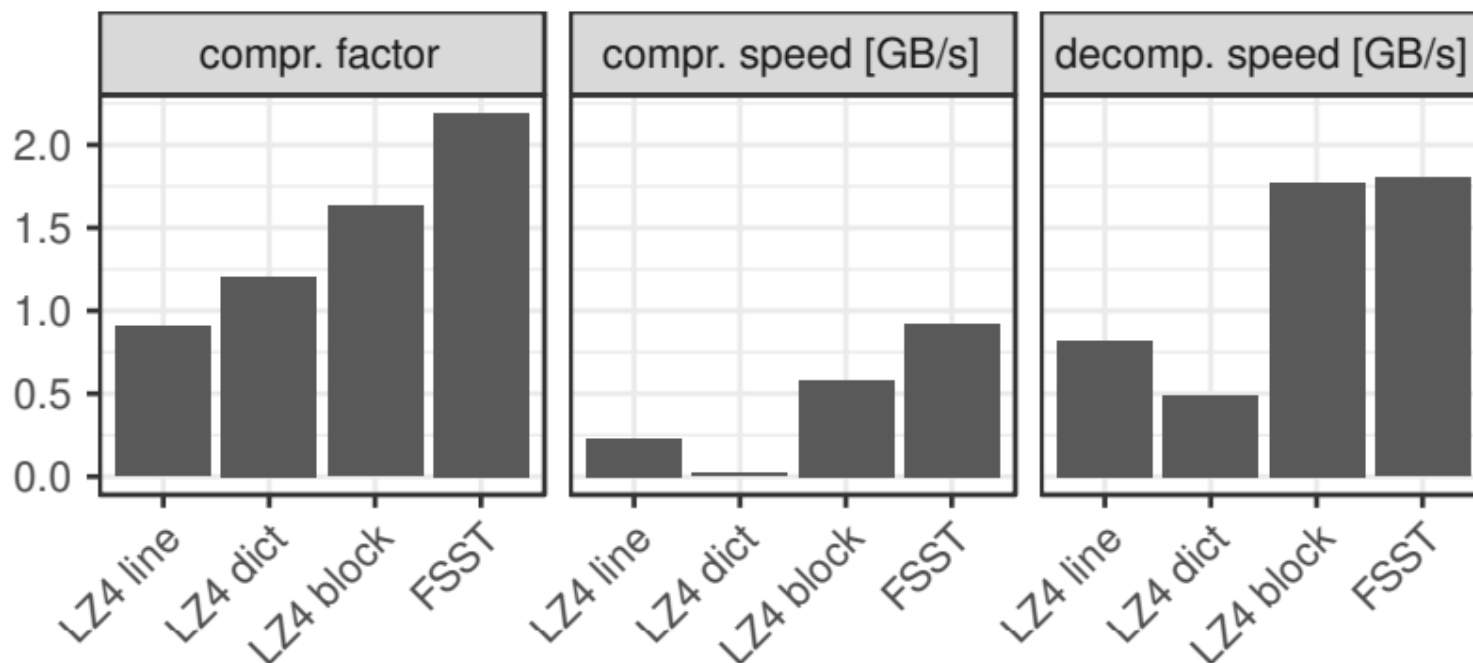


Figure 4: With LZ4 short strings do not compress well, even with a pre-generated dictionary.

Random Access: FSST vs LZ4

- Use case: **Scan** with a **pushed down predicate** (selection % on X axis)
 - LZ4 must decompress all strings, FSST only the selected tuples
 - FSST might even choose not to decompress strings (would even be faster)

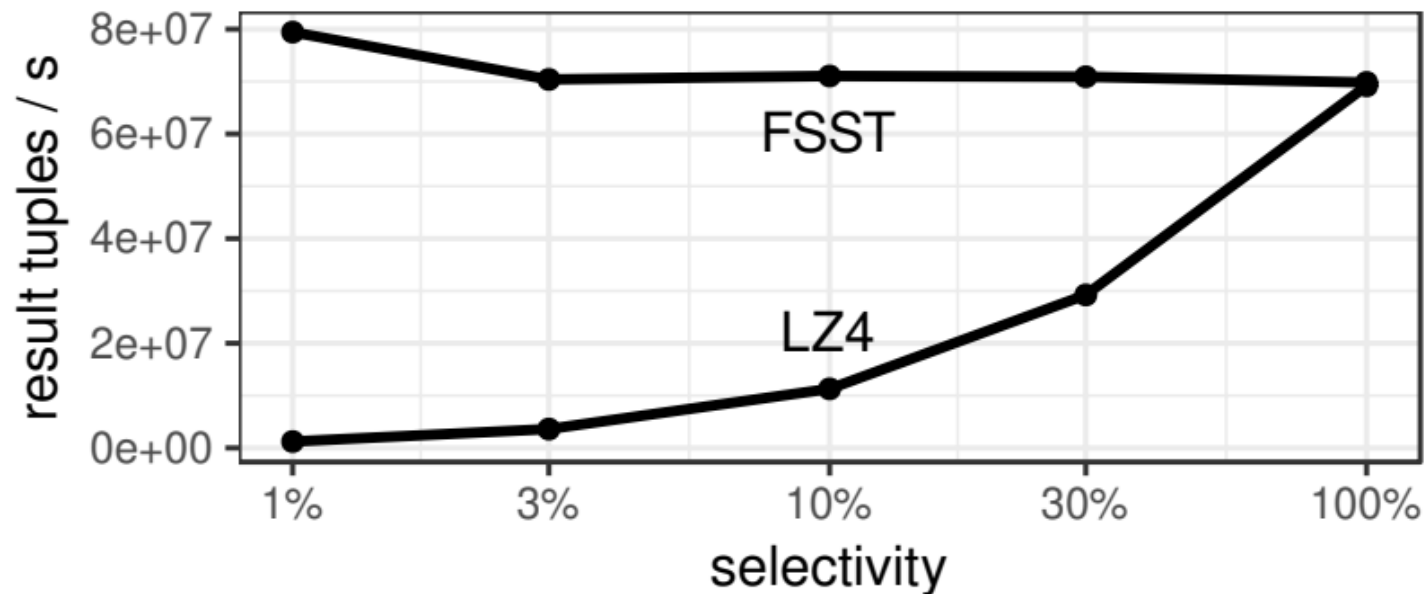


Figure 5: Selective queries are fast in FSST due to random access to individual values.

SIMD vs Scalar

- Encoding performance
 - 400MB/s scalar
 - 900MB/s AVX512
- Note:
 - Without AVX512, LZ4 compression is faster than FSST
 - But FSST speed still 70% of LZ4

<i>simd₁</i>	<i>simd₂</i>	<i>simd₃</i>	<i>simd₄</i>	<i>adptv</i>	
8.01	5.36	4.98	5.16	11.26	firstname
7.97	5.08	4.17	4.42	5.09	hex
8.70	5.51	5.07	4.69	10.82	city
5.29	3.57	2.85	3.26	9.00	genome
8.05	5.12	4.18	4.44	10.35	lastname
7.82	5.03	4.68	4.61	13.41	credentials
7.35	5.02	4.36	4.52	11.46	street
3.87	2.76	2.23	2.42	5.37	c_name
9.36	5.91	5.03	5.31	11.50	yago
9.10	5.74	4.82	5.14	11.36	movies
7.19	4.60	4.03	4.25	10.13	email
9.37	5.92	4.93	5.30	11.88	wiki
8.55	5.51	4.98	5.18	13.05	faust
5.59	3.63	3.16	3.42	9.78	l_comment
7.35	4.56	4.13	4.23	10.50	hamlet
6.36	4.08	3.40	3.59	8.06	uuid
5.64	3.70	3.05	3.14	5.73	location
7.50	4.72	4.03	4.28	10.12	urls2
6.32	4.04	3.51	3.74	8.55	urls
9.19	5.78	5.04	5.15	13.74	chinese
8.22	5.27	4.69	5.06	14.47	japanese
4.70	3.11	2.77	3.07	9.14	ps_comment
8.21	5.27	4.46	4.87	12.65	wikipedia
7.38	4.75	4.11	4.32	10.32	

Conclusion

- Databases are full of strings (see Public BI benchmark, DBtest “get real” paper)
 - String processing is a big bottleneck (CPU, RAM, network, disk)
 - **String compression** is therefore a good idea (less RAM, network, disk)
 - Operating on compressed strings is **very beneficial**
- FSST provides:
 - **random access to compressed strings!**
 - comparable/better (de)compression **speed** and **ratio** than the fastest general purpose compression schemes (LZ4)
- Useful opportunities of FSST:
 - Compressed execution, comparisons on compressed data
 - Late decompression (strings-stay-string). Has 0-terminated mode.
 - Easy integration in existing (database) systems
- **MIT licensed, code, paper + replication package** github.com/cwida/fsst